

Prompt Engineering Study Handbook

From Beginner to Advanced — Your Complete Reference Guide

Based on ABTalks × 60 Days of AI Prompt Engineering Roadmap

21

Chapters

50

Templates

30

Day Plan

TABLE OF CONTENTS

Ch 1	What Is Prompt Engineering?
Ch 2	Fundamentals of Prompt Engineering
Ch 3	The R-C-I-F-S Core Framework
Ch 4	Structured Prompting
Ch 5	Role-Based Prompting
Ch 6	Chain-of-Thought Prompting
Ch 7	Zero-Shot Prompting
Ch 8	Few-Shot Prompting
Ch 9	XML Prompting
Ch 10	Context Engineering
Ch 11	Prompt Chaining
Ch 12	AI Agents & Agentic Prompting
Ch 13	Claude Best Practices
Ch 14	Advanced Prompting Techniques
Ch 15	Business Use Cases
Ch 16	Core Prompt Templates Library
Ch 17	Cheat Sheets & Summary Tables
Ch 18	Interview Preparation Q&A;
Ch 19	50 Most Useful Prompt Templates
Ch 20	30-Day Practice Plan
Ch 21	Beginner-to-Advanced Learning Path

60 Days -
60 Days Claude Challenge
60 Days Claude Challenge by ABTalks
60 Days Claude Challenge by ABTalks
60 Days Claude Challenge by ABTalks
60 Days Claude Challenge by ABTalks

CHAPTER 1

What Is Prompt Engineering?

1.1 Definition

Prompt engineering is the practice of designing, structuring, and refining instructions (prompts) that guide an AI language model to produce accurate, relevant, and high-quality outputs — reliably, on the first try.

1.2 Why It Matters

- Saves time by reducing back-and-forth corrections
- Unlocks capabilities most users never discover
- Essential for freelancing, consulting, product development, and enterprise AI
- Differentiates casual users from professional AI practitioners
- Applicable across writing, coding, analysis, design, education, marketing

1.3 The Quality Ladder

Level	Example Prompt	Strength / Problem
Weak	write about marketing	Vague — generic essay
Okay	write 5 marketing tips for a small bakery	Better, but lacks persona/format
Good	Act as a marketing consultant. 5 low-budget Instagram posts. Bullets. format	Better. Bullets. format
Pro	Act as marketing consultant, 10 yrs food retail. Budget: \$1k. 5 ideas + roleplay action, outcome. Table.	5k. 5 ideas + roleplay action, outcome. Table.

■ You don't need fancy words — you need specificity. Every detail removes a guess the AI would otherwise make.

1.4 Common Beginner Mistakes

- Writing one vague line and hoping for the best
- Cramming multiple unrelated tasks into a single prompt
- Assuming the AI already knows your context
- Accepting the first draft as final
- Forgetting to specify who the output is for

CHAPTER 2

Fundamentals of Prompt Engineering

2.1 How AI Models Read a Prompt

AI models process your entire prompt as tokens (chunks of text), predicting the most statistically likely continuation. This means:

- **Order matters** — context given early shapes what comes later
- **Clarity matters** — ambiguity leads to averaged, generic outputs
- **Structure matters** — well-organised prompts produce well-organised responses

2.2 Prompt Types Overview

Type	Description	Best For
Zero-Shot	No examples; AI uses training alone	Quick tasks, factual queries
Few-Shot	1–3 examples provided in the prompt	Tone/style matching, structured outputs
Chain-of-Thought	Ask AI to reason step by step	Math, logic, complex decisions
Role-Based	Assign AI a specific expert persona	Expert-level output in any domain
XML Structured	Use structured markup tags	API workflows, agentic systems
Prompt Chaining	Break task into sequential prompts	Long docs, multi-step processes

2.3 Key Vocabulary

Term	Meaning
Token	A chunk of text (~4 chars or ¾ of a word)
Context window	Maximum tokens the AI can process in one session
Temperature	Controls creativity/randomness (0 = precise, 1 = creative)
System prompt	Background instruction given before the conversation
Iteration	Refining a prompt based on the output received
Hallucination	When AI produces confident but incorrect information

2.4 Pro Tips: Foundations

- *Specificity beats cleverness. Be direct, not poetic.*
- *Treat your first prompt as a draft. Expect to iterate 2–3 times.*
- *State the audience, goal, and format in every serious prompt.*

CHAPTER 3

The R-C-I-F-S Core Framework

The most important structure in this handbook. Master this and you will consistently produce professional-quality prompts.

3.1 The Five Elements

Element	What It Is	Example
R — Role	Expert persona the AI should adopt	Act as a certified financial planner with 15 years of experience
C — Context	Background: audience, goal, constraints	Audience: first-year college students; goal: start SIP with ₹500/month
I — Instruction	The single, clear task you want	Create a 3-step beginner investing plan
F — Format	How output should look	3-column table: Step Action Expected Outcome
S — Structure	Logical order of the response	Start with summary, then plan, end with one caution

3.2 The Copy-Paste Skeleton

Role: You are a [profession] with [X] years of experience in [domain].
Context: I need this for [audience] because [goal].
Constraints: [tone / length / must-include / avoid].
Instruction: [The one clear task you want completed].
Format: [bullets / table / 150-word paragraph / numbered list].
Structure: Start with [X], then cover [Y], end with [Z].

3.3 Do's and Don'ts

■ Do This	■ Avoid This
Be specific about audience and goal	Write one vague line and hope
Give an example of what "good" looks like	Cram five unrelated tasks in one prompt
Ask for one thing at a time	Assume AI knows your context
Specify length and format	Accept the first draft as final
Iterate — refine instead of restarting	Forget to say who it's for

3.4 Full Example Prompt

Role: You are a professional content strategist with 8 years of experience writing for LinkedIn thought leaders.
Context: My audience is mid-career professionals (28-40 years old) who feel too busy to read. Goal: inspire them to start reading 10 minutes a day.
Instruction: Write a LinkedIn post about the ROI of daily reading.
Format: 150-180 words. Hook line, 3 bullet points, closing CTA question.
Structure: Bold stat or surprising hook -> 3 benefit bullets -> 1-line CTA.

CHAPTER 4

Structured Prompting

Structured prompting organises your prompt using deliberate formatting — labels, separators, sections, and clear delimiters — so the AI can parse your intent without ambiguity.

4.1 Techniques

Labelled Sections

```
ROLE: Senior UX designer
TASK: Review this app flow
INPUT: [paste the flow here]
OUTPUT FORMAT: Bullet list of issues and one-line fixes
```

Delimiters

```
Summarise the text between the triple dashes.
---
[Your long text here]
---
Keep summary under 100 words.
```

4.2 When to Use Structured Prompting

Situation	Benefit
Multi-part instructions	Clear separation prevents confusion
Pasting user content	Delimiters isolate data from instructions
Specifying output format	Labels make format requirements explicit
API / developer workflows	Machine-readable, consistent parsing

4.3 Common Mistakes

- Mixing instruction and input without clear separation
- Giving format instructions after the content without a clear label
- Using inconsistent markers across a long prompt

CHAPTER 5

Role-Based Prompting

Role-based prompting assigns the AI a specific expert persona before asking it to complete a task. The persona shapes vocabulary, tone, depth, assumptions, and framing.

5.1 Role Formula

```
Act as a [title] with [X] years of experience in [industry/domain],  
specialising in [specific area], working with [type of clients/audience].
```

5.2 Simple vs Enhanced Roles

Simple Role	Enhanced Role
Act as a doctor	Act as a general physician with 12 years of practice in urban India, experienced in explaining medical concepts to first-generation health seekers
Act as a writer	Act as a senior copywriter at a D2C brand, specialising in conversion-focused landing page copy for health and wellness products
Act as a teacher	Act as a Class 10 Maths teacher skilled at breaking down concepts for students with limited prior foundation

5.3 Persona Layering Checklist

- **Title** — What is their job?
 - **Experience** — How many years?
 - **Industry** — Which sector?
 - **Audience they serve** — Who do they work with?
 - **Communication style** — Formal? Conversational? Technical?
- Try the same question across 3 different expert personas and compare. The contrast shows how powerfully roles shape responses.

CHAPTER 6

Chain-of-Thought (CoT) Prompting

CoT prompting asks the AI to reason through a problem step by step before delivering a final answer — dramatically improving accuracy on complex tasks.

6.1 Activation Phrases

- "Think step by step"
- "Walk me through your reasoning"
- "Break this down into steps before answering"
- "Reason through this before giving a final answer"

6.2 Standard CoT Template

```
I need to solve: [your problem].
Think step by step:
1. List what we know
2. Break the problem into parts
3. Solve each part
4. Give a clearly labelled FINAL ANSWER
```

6.3 Zero-Shot CoT vs Few-Shot CoT

Type	How It Works
Zero-Shot CoT	Add "Think step by step" to any prompt — no examples needed
Few-Shot CoT	Provide 1–2 worked examples with visible reasoning before the real question

6.4 Best Use Cases

- Financial calculations and budgeting
- Legal reasoning and case analysis
- Technical debugging
- Strategic decisions and trade-off comparisons
- Project planning and scheduling

CHAPTER 7

Zero-Shot Prompting

Zero-shot prompting gives the AI a task with no examples — relying entirely on the model's pre-trained knowledge and your instruction clarity.

7.1 When It Works Best

- Clear, well-defined tasks (summarise, translate, explain, list)
- Tasks where quality standards are flexible
- Quick first-draft generation
- Common formats (emails, blog intros, definitions)

7.2 Comparison: Basic vs Improved Zero-Shot

Version	Prompt
Basic Zero-Shot	Explain blockchain in 3 sentences for a 16-year-old.
Improved (Role + Format)	Act as a high school science teacher. Explain blockchain in exactly 3 sentences. Avoid jargon. End with one real-world analogy.

7.3 Limitations

- Style-specific outputs may be inconsistent without examples
- Niche or domain-specific tasks benefit from few-shot instead
- Complex or multi-step tasks need CoT or chaining

CHAPTER 8

Few-Shot Prompting

Few-shot prompting provides 1–3 examples of the exact input-output pattern you want, so the AI learns your style, tone, and structure and replicates it.

8.1 The Formula

```
[Optional: brief task description]

Example 1:
Input: [sample input]
Output: [desired output]

Example 2:
Input: [sample input]
Output: [desired output]

Now complete:
Input: [your actual input]
Output:
```

8.2 Example — LinkedIn Hooks

```
You write punchy LinkedIn hooks for B2B founders. Match this style:

Example 1:
Topic: Email marketing
Output: "Most founders ignore email. Their competitors don't."

Example 2:
Topic: Cold outreach
Output: "Your cold email isn't failing because of the subject line. It's the first sentence."

Now write a hook for:
Topic: Hiring your first salesperson
Output:
```

8.3 Few-Shot vs Zero-Shot Comparison

Dimension	Zero-Shot	Few-Shot
Examples needed	None	1–3
Style control	Low	High
Best for	Common formats	Niche/style-specific outputs
Sweet spot	—	2 examples

■ 2 examples is usually the sweet spot. One may not generalise; three or more can over-constrain. Always use your best work as examples.

CHAPTER 9

XML Prompting

XML prompting uses structured XML tags to clearly separate instructions, data, context, and output specifications with machine-readable precision. Heavily recommended for API and agentic workflows, and especially effective with Claude.

9.1 Basic Structure

```
Summarise the customer review below

Product review for an e-commerce platform.
Audience: operations team reviewing product quality.

[Paste the customer review here]

- Sentiment: Positive / Negative / Neutral
- Key Issue (if any): one sentence
- Suggested Action: one sentence
```

9.2 When to Use XML

Use Case	Why XML Helps
Document analysis pipelines	Clearly separates document from instruction
Multi-step workflows	Each step can be tagged and referenced
Extracting structured data	Output tags enable reliable parsing
System prompts for products	Clean separation of persona, rules, and task

9.3 Claude-Specific Advantage

Claude responds exceptionally well to XML-structured prompts. When building production workflows or system prompts, XML is the preferred format for clarity and reliability.

CHAPTER 10

Context Engineering

Context engineering is the practice of managing what information the AI has access to — including what to include, what to exclude, how to compress it, and when to reset it.

10.1 Context Window Reference

Model Type	Approximate Context Window
Standard models	8,000 – 32,000 tokens
Extended models	100,000 – 200,000 tokens
Claude 3+	Up to 200,000 tokens

10.2 The Memory Block Technique

When a conversation grows long, summarise key context into a compact block and re-paste it at the start of a new session:

```
MEMORY BLOCK (paste at start of each new session):
- Project: Redesign of loyalty program
- My role: UX Consultant
- Client constraints: Budget under Rs.10L, launch by March
- Decisions made: Tier system confirmed (Silver/Gold/Platinum)
- Last task completed: User journey map draft
- Next task: Wireframe for onboarding flow
```

10.3 The Caveman Trick (Context Compression)

Normal	Compressed (Caveman)
We have decided that the app will target urban women aged 22–35 who are interested in sustainable fashion and income.	Target: urban women, 22-35, sustainable fashion, high income.

10.4 Fresh vs Continue Decision


Situation	Recommendation
New unrelated task	Start fresh conversation
Continuing same project	Keep thread, add memory block if long
AI seems confused or off-track	Fresh session with clean compressed summary
Output quality degrading	Fresh session with better initial prompt

CHAPTER 11

Prompt Chaining

Prompt chaining breaks a complex task into a sequence of smaller, focused prompts — where the output of one prompt becomes the input to the next.

11.1 Chain Design Pattern



11.2 Example — Blog Post Chain

Step	Prompt
1 — Research	Act as a content strategist. List 7 data-backed reasons why sleep deprivation harms productivity. One statistic per point. Bullet list.
2 — Outline	Using these 7 points [paste Output 1], create a structured outline for a 1,000-word blog post targeting stressed corporate professionals.
3 — Draft	Using this outline [paste Output 2], write the full 1,000-word blog post. Tone: conversational and empathetic.
4 — Polish	Review this blog post [paste Output 3]: 1. Improve headline 2. Tighten paragraphs 3. Strengthen CTA.

11.3 Use Cases

- Long-form content (books, reports, courses)
- Research → analysis → recommendation workflows
- Data cleaning → analysis → visualisation scripts
- Job application (research → tailor CV → write cover letter)
- Product development (ideation → user stories → feature specs)

CHAPTER 12

AI Agents & Agentic Prompting

An AI agent takes autonomous sequences of actions — using tools, making decisions, and operating across multiple steps — to complete a complex goal with minimal human intervention.

12.1 Key Concepts

Concept	Description
Tool use	AI can call external functions (web search, code execution, APIs)
Memory	Persistent context across sessions (short-term or long-term)
Planning	AI breaks a goal into sub-tasks autonomously
Action loops	AI takes action, observes result, takes next action
Human-in-the-loop	Points where human review or approval is required

12.2 System Prompt Template for Agents

```
You are a research assistant agent.
Goal: answer questions by searching, synthesising, and citing sources.

TOOLS: web_search, document_reader, calculator

RULES:
- Verify claims using at least 2 sources
- Ask for clarification before irreversible actions
- Summarise your plan before executing
- Report errors immediately

OUTPUT: 1. Action taken 2. Finding 3. Source 4. Confidence: High/Medium/Low
```

12.3 Safety Principles

- **Minimal footprint:** Request only necessary permissions
- **Confirm before irreversible actions:** File deletion, sending emails, payments
- **Pause points:** Build in check-ins at key decision points
- **Clear error handling:** Tell the agent what to do with unexpected input

CHAPTER 13

Claude Best Practices

13.1 What Makes Claude Different

Feature	Detail
Instruction-following	Strong accuracy across complex, multi-part tasks
Context window	Up to 200K tokens — ideal for long document analysis
Nuanced reasoning	High accuracy on analytical and sensitive tasks
Structured prompts	Excels with XML and labelled-section formats
Careful responses	Thoughtful handling of ambiguous or sensitive areas

13.2 Claude-Specific Tips

Tip	How To Apply
Use XML tags	Claude responds exceptionally well to XML-structured prompts
Be explicit about thinking	Use: "Think carefully before answering" or "Work through this step by step"
Leverage system prompt	Put all persona, rules, constraints in system prompt via API — keep user messages clean
Long-context strength	Feed entire contracts, reports, or transcripts for deep analysis
Self-critique loop	Add: "Review your answer. What is the weakest part? How would you improve it?"

13.3 Claude System Prompt Template (API)

```
# SYSTEM PROMPT
You are [role] with [experience].
Your communication style is [tone].
Always [key rule 1] and [key rule 2].
Never [key restriction].

# USER MESSAGE
[Clear single instruction]
[Background, audience, goal]
[User content or question]
[Exact format required]
```

CHAPTER 14

Advanced Prompting Techniques

14.1 Self-Critique Prompting

Ask the AI to review and improve its own output.

```
Review your previous response.  
Identify the 2 weakest sections.  
Rewrite them to be 30% stronger.
```

14.2 Persona Contrast Prompting

Ask the same question from multiple expert angles for richer insights.

```
Answer from 3 perspectives:  
1. A risk-averse CFO  
2. An aggressive growth-stage VC  
3. A bootstrapped founder  
Question: Should we raise funding now or wait 12 months?
```

14.3 The Socratic Prompt

Instead of asking for an answer, ask the AI to ask you the right questions first.

```
I want to build a personal brand on LinkedIn.  
Don't give me a strategy yet.  
Ask me the 5 most important questions you need answered to give me tailored advice.
```

14.4 Reverse Prompting

Ask the AI to tell you what a great prompt for your task would look like.

```
I want to write high-converting product descriptions for an e-commerce site.  
What would the ideal prompt for this task look like?  
Give me a ready-to-use template.
```

14.5 Devil's Advocate

Get the strongest counter-arguments before committing to a decision.

```
I've decided to [your decision].  
Play devil's advocate: give me the 5 strongest arguments AGAINST this decision.  
Be rigorous and don't soften the critique.
```

14.6 Constraint-Based Creativity

Adding constraints improves creative output.

```
Write a tagline for a sustainable water bottle brand.  
Constraints: Exactly 5 words. Must include a verb.  
Cannot use: eco, green, sustainable, planet, earth.
```


CHAPTER 15

Business Use Cases

Marketing & Content

Task	Prompting Approach
Social media captions	Few-shot with style examples
Blog posts	Prompt chaining (research → outline → draft → edit)
Ad copy	Role-based (senior copywriter) + constraint-based
Email campaigns	R-C-I-F-S with audience segmentation in context
SEO content	Role + Chain-of-thought for keyword integration

Sales & Business Development

Task	Prompting Approach
Cold email sequences	Few-shot with your best-performing email as example
Pitch decks	Chaining (research → narrative → slides outline)
Proposal writing	Role-based (proposal writer) + context (client brief)
Objection handling	Persona contrast (buyer's perspective)
Follow-up messages	Template injection with dynamic context

Operations & Productivity

Task	Prompting Approach
Meeting summaries	Zero-shot with structured output format
SOPs	Chain-of-thought + numbered steps format
Email triage	XML classification prompt
Hiring (JDs, interviews)	Role-based (HR professional)
Report writing	Structured prompting with section scaffolding

CHAPTER 16

Core Prompt Templates Library

The Explainer

Explain [TOPIC] to me like I'm a [beginner / busy manager].
Use a simple analogy, then give 3 key takeaways as bullets. Keep it under 150 words.

The Professional Writer

Act as a [ROLE]. Write a [email / message / post] to [RECIPIENT].
Context: [what happened / what I need].
Tone: [warm but professional]. Length: under [120] words.
End with one clear call to action.

The Step-by-Step Solver

I need to solve: [YOUR PROBLEM].
Think step by step: list what we know, break it into parts,
solve each part, then give a clearly labelled FINAL ANSWER.

The Rewriter

Rewrite the text below to be [shorter / more formal / more persuasive].
Keep the core message intact. Target length: [word count].
Audience: [who will read this].
Original text: [PASTE TEXT]

The Brainstorm

Act as a [creative strategist / product manager / marketing consultant].
Give me 10 ideas for [GOAL/CHALLENGE].
Audience: [TARGET]. Constraint: [budget / time / channel].
Format: Numbered list. One sentence per idea.

CHAPTER 17

Cheat Sheets & Summary Tables

17.1 Technique Selection Guide

Your Goal	Best Technique
Quick simple task	Zero-Shot
Match a specific style	Few-Shot
Complex reasoning	Chain-of-Thought
Expert-level output	Role-Based
Long complex project	Prompt Chaining
API / developer use	XML Structured
Long documents	Context Engineering
Multiple perspectives	Persona Contrast
Idea generation	Constraint-Based Creativity
Improve AI output	Self-Critique Prompting

17.2 Prompt Quality Checklist

- Is there a clear role assigned?
- Is the audience specified?
- Is the goal / purpose stated?
- Is the task a single clear instruction?
- Is the format specified (length, structure)?
- Are constraints mentioned?
- Have I provided an example if needed?
- Am I prepared to iterate on the first output?

17.3 Prompt Failure Diagnosis

Problem	Likely Cause	Fix
Too generic	No role or audience	Add role + audience context
Wrong format	No format instruction	Specify bullets/table/word count
Too long	No length constraint	Add "under X words"
Off-topic	Task too vague	Use one clear verb + object
Missing my style	No examples	Add 1–2 few-shot examples
Incorrect reasoning	No step-by-step	Add "think step by step"
Loses thread	Context window issue	Add memory block

17.4 The 60-Second Prompt Formula

"Act as a [ROLE] with [EXPERIENCE].
For [AUDIENCE], [TASK VERB] [OBJECT].
Context: [GOAL + CONSTRAINT].
Format: [OUTPUT FORMAT].
Length: [WORD COUNT]."

Desired Effect	Phrase to Use
Style control	"Match this style: [example]"
Conciseness	"Under [X] words. Be direct."
Self-improvement	"Critique your answer and improve it"
Multiple views	"Answer from 3 perspectives: [A], [B], [C]"
Deeper questions	"Don't answer yet. Ask me the 5 most important questions first."

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

60 Days Claude Challenge by ABTalks

CHAPTER 18

Interview Preparation Q&A;

Master these 10 questions to ace any prompt engineering interview or client conversation.

Q1: What is prompt engineering and why is it important?

Prompt engineering is the skill of crafting precise instructions for AI models to produce reliable, high-quality outputs. It matters because the same AI can deliver vastly different results depending on how it is instructed — making it valuable across content creation, software development, data analysis, and business automation.

Q2: Explain the R-C-I-F-S framework.

R-C-I-F-S stands for Role, Context, Instruction, Format, and Structure. Role sets the AI's expert persona; Context gives background and constraints; Instruction states the single clear task; Format specifies the output presentation; Structure defines the logical order of the response.

Q3: What is Chain-of-Thought prompting and when do you use it?

CoT prompting asks the AI to reason step by step before giving a final answer. Most useful for complex, multi-step problems — mathematical calculations, logical reasoning, troubleshooting, and decision comparisons — where jumping to a conclusion increases error risk.

Q4: What is the difference between zero-shot and few-shot prompting?

Zero-shot gives the AI no examples — it relies on training alone. Few-shot provides 1–3 input-output examples so the AI learns the exact style required. Few-shot is more reliable for style-specific, format-sensitive, or niche outputs.

Q5: What is a context window and how does it affect prompting?

The context window is the maximum text an AI can process in one interaction. As conversations grow longer, earlier content receives less attention. Good context engineering means summarising earlier content into memory blocks and knowing when to start a fresh session.

Q6: What is prompt chaining?

Breaking a complex task into a sequence of smaller prompts, where each output feeds into the next. Enables higher quality on tasks too complex for a single prompt, mirroring a professional workflow: research, outline, draft, edit.

Q7: What are XML prompts and why are they used in production?

XML prompts use structured tags to separate instructions, data, context, and output requirements. Preferred in production because they reduce ambiguity, enable reliable output parsing, and keep complex multi-part instructions organised.

Q8: How do you handle a prompt that keeps producing poor results?

Use Prompt Failure Diagnosis: identify whether the issue is vagueness (add role/audience), format (specify explicitly), length (add word count), reasoning (add step-by-step), or style (add few-shot examples). Iterate with targeted refinements rather than rewriting the whole prompt.

Q9: What is agentic prompting?

Writing instructions for AI systems that take autonomous action sequences using tools, memory, and planning. Includes specifying available tools, behaviour rules, decision checkpoints, and error-handling protocols to ensure the agent acts safely and predictably.

Q10: How would you build a prompt library for a business?

Identify the 10–15 most repeated tasks. Write and test an R-C-I-F-S prompt for each. Save as templates with [BLANK] placeholders. Organise by department or use case. Assign ownership, version control, and a review cycle to keep them updated as AI models evolve.

CHAPTER 19

50 Most Useful Prompt Templates

Copy these templates, fill the brackets, and you are immediately prompting better than most people.

Writing & Content

1. Blog Post Generator

Act as a senior content writer. Write a [word count]-word blog post for [audience] about [topic]. Use a hook, 3-4 subheadings, and a CTA at the end. Tone: [tone].

2. LinkedIn Post

Act as a LinkedIn thought leader in [industry].
Write a 150-word post about [insight/lesson].
Start with a provocative one-line hook. End with a question to drive comments.

3. Email Subject Lines

Write 10 email subject lines for a campaign about [offer/topic].
Audience: [target]. Goal: [open rate/clicks/conversions].
Vary between curiosity, benefit, and urgency styles.

4. Product Description

Act as a D2C copywriter. Write a 100-word product description for [product].
Audience: [buyer persona]. Highlight: [top 3 benefits]. Tone: [aspirational/practical/playful].

5. Cold Email

Write a cold outreach email to [role] at [company type]. My offer: [what I provide].
Their pain: [problem I solve]. Length: under 120 words. End with low-friction CTA.

Analysis & Research

6. Competitor Analysis

Act as a business analyst. Compare [Company A] and [Company B] on: pricing, target market, key features, strengths, weaknesses. Format: comparison table + 2-sentence summary.

7. SWOT Analysis

Conduct a SWOT analysis for [business/idea].
Format: 4-quadrant table. 3 points per quadrant.
End with: "Top strategic priority based on this SWOT:"

8. Market Research Summary

Summarise key findings for [industry/market] in India in 2025. Include: market size, growth rate, top players, consumer trends, one emerging opportunity. Format: structured report.

9. Data Interpretation

Here is a dataset: [paste data]. Think step by step.
Identify: top trend, biggest anomaly, and one actionable insight.
Format: 3 bullets followed by a one-paragraph recommendation.

10. Literature Review

Summarise the key arguments from this text: [paste text].
Identify: main thesis, supporting evidence, limitations, relevance to [my topic].
Format: structured notes under 300 words.

Professional Communication

11. Performance Review

Act as an HR professional. Write a performance review for an employee who [achievements].
Tone: constructive and encouraging. Length: 200 words.
Include: 2 strengths, 1 growth area, specific example for each.

12. Meeting Agenda

14. Proposal Executive Summary

Write a 200-word executive summary for a proposal to [client type] for [service/project].
Highlight: problem, solution, expected outcome, and why us. Tone: confident and credible.

15. Feedback Message

Help me give constructive feedback to [role] about [issue].
Goal: [desired behaviour change]. Use SBI model: Situation → Behaviour → Impact.
Tone: supportive. Under 120 words.

Learning & Education

16. Concept Explainer

Explain [concept] to a [beginner / 12-year-old / professional in another field].
Use one simple analogy, then 3 key takeaways as bullets. Under 150 words.

17. Study Plan Generator

Create a 30-day study plan for learning [skill/subject]. I can dedicate [X hours/day].
Current level: [beginner/intermediate]. Goal: [specific outcome]. Format: week-by-week breakdown.

18. Quiz Generator

Create a 10-question multiple choice quiz on [topic].
Difficulty: [beginner/intermediate/advanced].
Format: question, 4 options (A-D), correct answer on next line.

19. Flashcard Creator

Create 15 flashcards for [topic].
Format: FRONT: [key term or question] | BACK: [definition or answer]
Keep each back under 20 words.

20. Socratic Learning

I want to learn about [topic]. Don't explain yet.
Ask me 5 questions to discover what I already know and what gaps I have.
Then give me a personalised explanation.

Strategy & Decision Making

21. Decision Framework

I'm deciding between [Option A] and [Option B]. Think step by step.
Evaluate both on: [criteria 1], [criteria 2], [criteria 3].
Format: comparison table, then recommendation + confidence level.

22. Devil's Advocate

I've decided to [decision]. Play devil's advocate.
Give me the 5 strongest arguments AGAINST this decision.
Be rigorous. Don't soften the critique.

23. Risk Assessment

Act as a risk consultant. Assess the top 5 risks in [project/decision].
For each: Likelihood (H/M/L), Impact (H/M/L), and one mitigation step. Format: table.

24. Goal Setting (SMART)

Help me set a 90-day goal for [objective]. Use the SMART framework.
Then break into 3 monthly milestones and 3 weekly habits. Format: structured plan.

25. Business Idea Validator

I have a business idea: [describe idea].
Evaluate on: problem clarity, target market, competition, monetisation, founder-market fit.
Score out of 10 for each with one-line reasoning. End with go/no-go recommendation.

Technical & Coding

26. Code Explainer

Explain this code to a junior developer: [paste code]

28. Documentation Writer

Write clear documentation for this function: [paste code]
Include: purpose, parameters, return value, edge cases, and one usage example.
Format: standard docstring style.

29. Architecture Advisor

Act as a senior software architect. I'm building [describe system].
Scale requirement: [users/requests]. Recommend a tech stack with justification.
Highlight the top 2 trade-offs.

30. Regex Generator

Write a regex pattern that matches [describe pattern].
Provide: the pattern, a plain English explanation,
and 3 example strings that would or would not match.

Freelancing & Consulting

31. Service Proposal

Write a project proposal for a [service] engagement with [client type].
Scope: [deliverables]. Timeline: [duration].
Format: Problem → Approach → Deliverables → Timeline → Investment.

32. Client Onboarding Email

Write a welcome email to a new client who has hired me for [service].
Include: what happens next, what I need from them, and timeline.
Tone: warm and professional. Under 200 words.

33. LinkedIn Profile Summary

Write a LinkedIn summary for a [role] with [years] years of experience in [skills].
Target audience: [type of client/employer]. Highlight: [top 3 achievements/skills].
Under 220 words. End with a clear value proposition line.

34. Scope of Work (SOW)

Create a Scope of Work document for [project type].
Include: objectives, deliverables, timeline, assumptions, out-of-scope items.
Format: professional document with headings.

35. Case Study Writer

Write a client case study for [project].
Structure: Client Background → Problem → Solution → Results → Testimonial placeholder.
Tone: credible and results-focused. Length: 300-400 words.

Productivity & Personal Use

36. Weekly Planner

Help me plan my week. Goals: [list goals]. Fixed commitments: [list].
Available work hours: [X/day].
Format: day-by-day schedule with time blocks. Prioritise by impact.

37. Journal Prompts

Give me 7 deep journal prompts for reflecting on [theme: growth/relationships/career].
Each prompt should be open-ended and encourage honest self-examination.

38. Habit Builder

I want to build the habit of [habit]. Current obstacle: [what stops me].
Design a 30-day habit plan using implementation intentions (When X happens, I will do Y).
Format: week-by-week table.

39. Book Summary

Summarise [book title] by [author].
Include: main thesis, 5 key ideas, best quote, and one actionable takeaway.
Format: structured notes. Under 300 words.

41. Persona Contrast

Answer this question from 3 distinct perspectives:

1. [Persona 1] 2. [Persona 2] 3. [Persona 3]

Question: [your question]. Format: one paragraph per perspective, labelled clearly.

42. Self-Critique

Review your previous answer.

Identify the 2 weakest parts and explain why they are weak.

Then rewrite those sections to be significantly stronger.

43. Reverse Engineering

Here is an excellent piece of [writing/code/strategy]: [paste example].

Analyse what makes it effective.

Extract a framework or template I can reuse for similar work.

44. Future Scenario Planner

Act as a futurist. Describe 3 scenarios for [topic/industry] in 5 years:

Optimistic, Realistic, Pessimistic.

For each: key drivers, key risks, one strategic recommendation. Format: 3-column table.

45. Teaching Assistant

I am going to teach [topic] to [audience]. Review my draft lesson plan: [paste plan].

Identify: gaps in explanation, concepts that need examples,
and the most likely point of confusion for the audience.

Specialised Use Cases

46. HR Screening Questions

Act as a senior hiring manager. Create 8 interview questions for a [role] position.

Mix: technical competency, behavioural (STAR format), and culture-fit.

Include what a strong answer would contain for each.

47. Press Release Writer

Write a press release for [announcement]. Company: [name]. Audience: [media type].

Format: Headline → Dateline → Lead → Quote → Body → Boilerplate.

Tone: professional and newsworthy.

48. SOP Generator

Write a Standard Operating Procedure for [process].

Audience: new team members with no prior experience.

Format: numbered steps with decision points. Include a "common mistakes" section at the end.

49. UX Copy Writer

Write microcopy for a [onboarding/checkout/error state] in a [type of app].

User emotion at this moment: [frustrated/excited/confused].

Tone: [warm/professional/playful]. Keep each string under 12 words.

50. Negotiation Prep

I'm preparing to negotiate [situation: salary/contract/price].

My position: [what I want]. Their likely position: [what they want].

Give: 3 opening lines, 5 anticipated objections with responses, and my BATNA.

CHAPTER 20

30-Day Practice Plan

Week 1: Foundation	
Day	Task
Day 1	Rewrite 3 vague prompts into specific ones. Compare outputs.
Day 2	Use the full R-C-I-F-S skeleton on a real work task.
Day 3	Try "think step by step" on a complex decision you face.
Day 4	Ask the same question as 3 different expert personas.
Day 5	Give 2 examples (few-shot) and compare to zero-shot output.
Day 6	Take one output and refine it with 3 follow-up nudges.
Day 7	Save your 3 best prompts as reusable templates.
Week 2: Intermediate	
Day	Task
Day 8	Write an XML-structured prompt for a complex task.
Day 9	Build a 3-prompt chain for a blog post or report.
Day 10	Use the Socratic prompt technique for a topic you want to learn.
Day 11	Use persona contrast on a business decision.
Day 12	Use the self-critique technique on your best recent AI output.
Day 13	Summarise a long document into a compact memory block.
Day 14	Review your prompt library. Improve 3 templates.
Week 3: Application	
Day	Task
Day 15	Write 5 prompts for your most common work tasks.
Day 16	Use constraint-based creativity for a marketing task.
Day 17	Use reverse prompting — ask AI what a great prompt looks like.
Day 18	Build a prompt for a business use case (sales, ops, HR).
Day 19	Use the devil's advocate technique on a real decision.
Day 20	Practice a few-shot data extraction prompt.
Day 21	Full project using chaining: research → outline → draft → edit.
Week 4: Mastery	
Day	Task
Day 22	Build a Claude system prompt for a persona you use regularly.
Day 23	Design a simple agentic prompt with tool use specifications.
Day 24	Create a prompt template library with 20 entries.
Day 25	Run the Chapter 18 interview Q&A from memory.

CHAPTER 21

Beginner-to-Advanced Learning Path

Stage 1: FOUNDATION — Beginner	Days 1–7	Get comfortable. Learn that specificity beats cleverness.
- Understand what a prompt is and how the AI reads it - Practice turning vague asks into specific ones - Always state the audience, goal, and desired format - Learn to give one clear instruction per prompt ■ Milestone: Write a prompt that gets a useful output on the first try		
Stage 2: STRUCTURE — Beginner+	Days 8–14	Use R-C-I-F-S deliberately on every important prompt.
- Add a role to shift the AI's expertise and tone - Layer in context: who, why, and constraints - Control the output format explicitly (tables, bullets, word counts) - Reorder elements and observe how answers change ■ Milestone: Save 5 reusable R-C-I-F-S templates		
Stage 3: REASONING — Intermediate	Days 15–18	Get better answers on complex tasks by asking AI to think in steps.
- Add "think step by step" for math, logic, and planning - Break big problems into ordered sub-questions - Ask for reasoning first, then a labelled final answer - Use for decisions, comparisons, and troubleshooting ■ Milestone: Use CoT to solve a real business or personal decision		
Stage 4: PERSONAS — Intermediate	Days 19–22	Turn the AI into any expert and teach it by example.
- Assign detailed personas: role + experience + industry + audience - Use few-shot: give 1–3 examples of the output you want - Compare the same question across different expert personas - Lock in tone, style, and structure via examples ■ Milestone: Produce expert-level output in a domain where you're not an expert		
Stage 5: CONTROL — Intermediate	Days 23–26	Run long conversations and complex workflows without losing quality.
- Understand and manage the context window - Use memory blocks for long projects - Use XML-structured prompts for complex multi-part tasks - Use prompt chaining for multi-step workflows ■ Milestone: Complete a long-form project using chaining		
Stage 6: MASTERY — Advanced	Days 27–30	Treat prompting as a loop, build assets, and think at systems level.
- Refine with follow-ups and self-critique loops - Save best prompts as templates with fill-in-the-blank slots - Build a personal prompt library for weekly recurring tasks - Design agentic prompts and system prompts for products ■ Milestone: Own a prompt library of 30+ reusable, tested templates		

Prompt Engineering Study Handbook

From Beginner to Advanced — Your Complete Reference Guide

Based on ABTalks × 60 Days of AI Prompt Engineering Roadmap

21

Chapters

50

Templates

30

Day Plan

TABLE OF CONTENTS

Ch 1	What Is Prompt Engineering?
Ch 2	Fundamentals of Prompt Engineering
Ch 3	The R-C-I-F-S Core Framework
Ch 4	Structured Prompting
Ch 5	Role-Based Prompting
Ch 6	Chain-of-Thought Prompting
Ch 7	Zero-Shot Prompting
Ch 8	Few-Shot Prompting
Ch 9	XML Prompting
Ch 10	Context Engineering
Ch 11	Prompt Chaining
Ch 12	AI Agents & Agentic Prompting
Ch 13	Claude Best Practices
Ch 14	Advanced Prompting Techniques
Ch 15	Business Use Cases
Ch 16	Core Prompt Templates Library
Ch 17	Cheat Sheets & Summary Tables
Ch 18	Interview Preparation Q&A;
Ch 19	50 Most Useful Prompt Templates
Ch 20	30-Day Practice Plan
Ch 21	Beginner-to-Advanced Learning Path

60 Days -
60 Days Claude Challenge
60 Days Claude Challenge by ABTalks
60 Days Claude Challenge by ABTalks
60 Days Claude Challenge by ABTalks
60 Days Claude Challenge by ABTalks

CHAPTER 1

What Is Prompt Engineering?

1.1 Definition

Prompt engineering is the practice of designing, structuring, and refining instructions (prompts) that guide an AI language model to produce accurate, relevant, and high-quality outputs — reliably, on the first try.

1.2 Why It Matters

- Saves time by reducing back-and-forth corrections
- Unlocks capabilities most users never discover
- Essential for freelancing, consulting, product development, and enterprise AI
- Differentiates casual users from professional AI practitioners
- Applicable across writing, coding, analysis, design, education, marketing

1.3 The Quality Ladder

Level	Example Prompt	Strength / Problem
Weak	write about marketing	Vague — generic essay
Okay	write 5 marketing tips for a small bakery	Better, but lacks persona/format
Good	Act as a marketing consultant. 5 low-budget Instagram posts. Bullets. format	Better. Bullets. format
Pro	Act as marketing consultant, 10 yrs food retail. Budget: \$1k. 5 ideas + roleplay action, outcome. Table.	5k. 5 ideas + roleplay action, outcome. Table.

■ You don't need fancy words — you need specificity. Every detail removes a guess the AI would otherwise make.

1.4 Common Beginner Mistakes

- Writing one vague line and hoping for the best
- Cramming multiple unrelated tasks into a single prompt
- Assuming the AI already knows your context
- Accepting the first draft as final
- Forgetting to specify who the output is for

CHAPTER 2

Fundamentals of Prompt Engineering

2.1 How AI Models Read a Prompt

AI models process your entire prompt as tokens (chunks of text), predicting the most statistically likely continuation. This means:

- **Order matters** — context given early shapes what comes later
- **Clarity matters** — ambiguity leads to averaged, generic outputs
- **Structure matters** — well-organised prompts produce well-organised responses

2.2 Prompt Types Overview

Type	Description	Best For
Zero-Shot	No examples; AI uses training alone	Quick tasks, factual queries
Few-Shot	1–3 examples provided in the prompt	Tone/style matching, structured outputs
Chain-of-Thought	Ask AI to reason step by step	Math, logic, complex decisions
Role-Based	Assign AI a specific expert persona	Expert-level output in any domain
XML Structured	Use structured markup tags	API workflows, agentic systems
Prompt Chaining	Break task into sequential prompts	Long docs, multi-step processes

2.3 Key Vocabulary

Term	Meaning
Token	A chunk of text (~4 chars or ¾ of a word)
Context window	Maximum tokens the AI can process in one session
Temperature	Controls creativity/randomness (0 = precise, 1 = creative)
System prompt	Background instruction given before the conversation
Iteration	Refining a prompt based on the output received
Hallucination	When AI produces confident but incorrect information

2.4 Pro Tips: Foundations

- **Specificity beats cleverness. Be direct, not poetic.**
- **Treat your first prompt as a draft. Expect to iterate 2–3 times.**
- **State the audience, goal, and format in every serious prompt.**

CHAPTER 3

The R-C-I-F-S Core Framework

The most important structure in this handbook. Master this and you will consistently produce professional-quality prompts.

3.1 The Five Elements

Element	What It Is	Example
R — Role	Expert persona the AI should adopt	Act as a certified financial planner with 15 years of experience
C — Context	Background: audience, goal, constraints	Audience: first-year college students; goal: start SIP with ₹500/month
I — Instruction	The single, clear task you want	Create a 3-step beginner investing plan
F — Format	How output should look	3-column table: Step Action Expected Outcome
S — Structure	Logical order of the response	Start with summary, then plan, end with one caution

3.2 The Copy-Paste Skeleton

Role: You are a [profession] with [X] years of experience in [domain].
Context: I need this for [audience] because [goal].
Constraints: [tone / length / must-include / avoid].
Instruction: [The one clear task you want completed].
Format: [bullets / table / 150-word paragraph / numbered list].
Structure: Start with [X], then cover [Y], end with [Z].

3.3 Do's and Don'ts

■ Do This	■ Avoid This
Be specific about audience and goal	Write one vague line and hope
Give an example of what "good" looks like	Cram five unrelated tasks in one prompt
Ask for one thing at a time	Assume AI knows your context
Specify length and format	Accept the first draft as final
Iterate — refine instead of restarting	Forget to say who it's for

3.4 Full Example Prompt

Role: You are a professional content strategist with 8 years of experience writing for LinkedIn thought leaders.
Context: My audience is mid-career professionals (28-40 years old) who feel too busy to read. Goal: inspire them to start reading 10 minutes a day.
Instruction: Write a LinkedIn post about the ROI of daily reading.
Format: 150-180 words. Hook line, 3 bullet points, closing CTA question.
Structure: Bold stat or surprising hook -> 3 benefit bullets -> 1-line CTA.

CHAPTER 4

Structured Prompting

Structured prompting organises your prompt using deliberate formatting — labels, separators, sections, and clear delimiters — so the AI can parse your intent without ambiguity.

4.1 Techniques

Labelled Sections

```
ROLE: Senior UX designer
TASK: Review this app flow
INPUT: [paste the flow here]
OUTPUT FORMAT: Bullet list of issues and one-line fixes
```

Delimiters

```
Summarise the text between the triple dashes.
---
[Your long text here]
---
Keep summary under 100 words.
```

4.2 When to Use Structured Prompting

Situation	Benefit
Multi-part instructions	Clear separation prevents confusion
Pasting user content	Delimiters isolate data from instructions
Specifying output format	Labels make format requirements explicit
API / developer workflows	Machine-readable, consistent parsing

4.3 Common Mistakes

- Mixing instruction and input without clear separation
- Giving format instructions after the content without a clear label
- Using inconsistent markers across a long prompt

CHAPTER 5

Role-Based Prompting

Role-based prompting assigns the AI a specific expert persona before asking it to complete a task. The persona shapes vocabulary, tone, depth, assumptions, and framing.

5.1 Role Formula

```
Act as a [title] with [X] years of experience in [industry/domain],  
specialising in [specific area], working with [type of clients/audience].
```

5.2 Simple vs Enhanced Roles

Simple Role	Enhanced Role
Act as a doctor	Act as a general physician with 12 years of practice in urban India, experienced in explaining medical concepts to first-generation health seekers
Act as a writer	Act as a senior copywriter at a D2C brand, specialising in conversion-focused landing page copy for health and wellness products
Act as a teacher	Act as a Class 10 Maths teacher skilled at breaking down concepts for students with limited prior foundation

5.3 Persona Layering Checklist

- **Title** — What is their job?
 - **Experience** — How many years?
 - **Industry** — Which sector?
 - **Audience they serve** — Who do they work with?
 - **Communication style** — Formal? Conversational? Technical?
- *Try the same question across 3 different expert personas and compare. The contrast shows how powerfully roles shape responses.*

CHAPTER 6

Chain-of-Thought (CoT) Prompting

CoT prompting asks the AI to reason through a problem step by step before delivering a final answer — dramatically improving accuracy on complex tasks.

6.1 Activation Phrases

- "Think step by step"
- "Walk me through your reasoning"
- "Break this down into steps before answering"
- "Reason through this before giving a final answer"

6.2 Standard CoT Template

```
I need to solve: [your problem].
Think step by step:
1. List what we know
2. Break the problem into parts
3. Solve each part
4. Give a clearly labelled FINAL ANSWER
```

6.3 Zero-Shot CoT vs Few-Shot CoT

Type	How It Works
Zero-Shot CoT	Add "Think step by step" to any prompt — no examples needed
Few-Shot CoT	Provide 1–2 worked examples with visible reasoning before the real question

6.4 Best Use Cases

- Financial calculations and budgeting
- Legal reasoning and case analysis
- Technical debugging
- Strategic decisions and trade-off comparisons
- Project planning and scheduling

CHAPTER 7

Zero-Shot Prompting

Zero-shot prompting gives the AI a task with no examples — relying entirely on the model's pre-trained knowledge and your instruction clarity.

7.1 When It Works Best

- Clear, well-defined tasks (summarise, translate, explain, list)
- Tasks where quality standards are flexible
- Quick first-draft generation
- Common formats (emails, blog intros, definitions)

7.2 Comparison: Basic vs Improved Zero-Shot

Version	Prompt
Basic Zero-Shot	Explain blockchain in 3 sentences for a 16-year-old.
Improved (Role + Format)	Act as a high school science teacher. Explain blockchain in exactly 3 sentences. Avoid jargon. End with one real-world analogy.

7.3 Limitations

- Style-specific outputs may be inconsistent without examples
- Niche or domain-specific tasks benefit from few-shot instead
- Complex or multi-step tasks need CoT or chaining

CHAPTER 8

Few-Shot Prompting

Few-shot prompting provides 1–3 examples of the exact input-output pattern you want, so the AI learns your style, tone, and structure and replicates it.

8.1 The Formula

```
[Optional: brief task description]

Example 1:
Input: [sample input]
Output: [desired output]

Example 2:
Input: [sample input]
Output: [desired output]

Now complete:
Input: [your actual input]
Output:
```

8.2 Example — LinkedIn Hooks

```
You write punchy LinkedIn hooks for B2B founders. Match this style:

Example 1:
Topic: Email marketing
Output: "Most founders ignore email. Their competitors don't."

Example 2:
Topic: Cold outreach
Output: "Your cold email isn't failing because of the subject line. It's the first sentence."

Now write a hook for:
Topic: Hiring your first salesperson
Output:
```

8.3 Few-Shot vs Zero-Shot Comparison

Dimension	Zero-Shot	Few-Shot
Examples needed	None	1–3
Style control	Low	High
Best for	Common formats	Niche/style-specific outputs
Sweet spot	—	2 examples

■ 2 examples is usually the sweet spot. One may not generalise; three or more can over-constrain. Always use your best work as examples.

CHAPTER 9

XML Prompting

XML prompting uses structured XML tags to clearly separate instructions, data, context, and output specifications with machine-readable precision. Heavily recommended for API and agentic workflows, and especially effective with Claude.

9.1 Basic Structure

```
Summarise the customer review below

Product review for an e-commerce platform.
Audience: operations team reviewing product quality.

[Paste the customer review here]

- Sentiment: Positive / Negative / Neutral
- Key Issue (if any): one sentence
- Suggested Action: one sentence
```

9.2 When to Use XML

Use Case	Why XML Helps
Document analysis pipelines	Clearly separates document from instruction
Multi-step workflows	Each step can be tagged and referenced
Extracting structured data	Output tags enable reliable parsing
System prompts for products	Clean separation of persona, rules, and task

9.3 Claude-Specific Advantage

Claude responds exceptionally well to XML-structured prompts. When building production workflows or system prompts, XML is the preferred format for clarity and reliability.

CHAPTER 10

Context Engineering

Context engineering is the practice of managing what information the AI has access to — including what to include, what to exclude, how to compress it, and when to reset it.

10.1 Context Window Reference

Model Type	Approximate Context Window
Standard models	8,000 – 32,000 tokens
Extended models	100,000 – 200,000 tokens
Claude 3+	Up to 200,000 tokens

10.2 The Memory Block Technique

When a conversation grows long, summarise key context into a compact block and re-paste it at the start of a new session:

```
MEMORY BLOCK (paste at start of each new session):
- Project: Redesign of loyalty program
- My role: UX Consultant
- Client constraints: Budget under Rs.10L, launch by March
- Decisions made: Tier system confirmed (Silver/Gold/Platinum)
- Last task completed: User journey map draft
- Next task: Wireframe for onboarding flow
```

10.3 The Caveman Trick (Context Compression)

Normal	Compressed (Caveman)
We have decided that the app will target urban women aged 22–35 who are interested in sustainable fashion and income.	Target: urban women, 22-35, sustainable fashion, high income.

10.4 Fresh vs Continue Decision


Situation	Recommendation
New unrelated task	Start fresh conversation
Continuing same project	Keep thread, add memory block if long
AI seems confused or off-track	Fresh session with clean compressed summary
Output quality degrading	Fresh session with better initial prompt

CHAPTER 11

Prompt Chaining

Prompt chaining breaks a complex task into a sequence of smaller, focused prompts — where the output of one prompt becomes the input to the next.

11.1 Chain Design Pattern



11.2 Example — Blog Post Chain

Step	Prompt
1 — Research	Act as a content strategist. List 7 data-backed reasons why sleep deprivation harms productivity. One statistic per point. Bullet list.
2 — Outline	Using these 7 points [paste Output 1], create a structured outline for a 1,000-word blog post targeting stressed corporate professionals.
3 — Draft	Using this outline [paste Output 2], write the full 1,000-word blog post. Tone: conversational and empathetic.
4 — Polish	Review this blog post [paste Output 3]: 1. Improve headline 2. Tighten paragraphs 3. Strengthen CTA.

11.3 Use Cases

- Long-form content (books, reports, courses)
- Research → analysis → recommendation workflows
- Data cleaning → analysis → visualisation scripts
- Job application (research → tailor CV → write cover letter)
- Product development (ideation → user stories → feature specs)

CHAPTER 12

AI Agents & Agentic Prompting

An AI agent takes autonomous sequences of actions — using tools, making decisions, and operating across multiple steps — to complete a complex goal with minimal human intervention.

12.1 Key Concepts

Concept	Description
Tool use	AI can call external functions (web search, code execution, APIs)
Memory	Persistent context across sessions (short-term or long-term)
Planning	AI breaks a goal into sub-tasks autonomously
Action loops	AI takes action, observes result, takes next action
Human-in-the-loop	Points where human review or approval is required

12.2 System Prompt Template for Agents

```
You are a research assistant agent.  
Goal: answer questions by searching, synthesising, and citing sources.  
  
TOOLS: web_search, document_reader, calculator  
  
RULES:  
- Verify claims using at least 2 sources  
- Ask for clarification before irreversible actions  
- Summarise your plan before executing  
- Report errors immediately  
  
OUTPUT: 1. Action taken 2. Finding 3. Source 4. Confidence: High/Medium/Low
```

12.3 Safety Principles

- **Minimal footprint:** Request only necessary permissions
- **Confirm before irreversible actions:** File deletion, sending emails, payments
- **Pause points:** Build in check-ins at key decision points
- **Clear error handling:** Tell the agent what to do with unexpected input

CHAPTER 13

Claude Best Practices

13.1 What Makes Claude Different

Feature	Detail
Instruction-following	Strong accuracy across complex, multi-part tasks
Context window	Up to 200K tokens — ideal for long document analysis
Nuanced reasoning	High accuracy on analytical and sensitive tasks
Structured prompts	Excels with XML and labelled-section formats
Careful responses	Thoughtful handling of ambiguous or sensitive areas

13.2 Claude-Specific Tips

Tip	How To Apply
Use XML tags	Claude responds exceptionally well to XML-structured prompts
Be explicit about thinking	Use: "Think carefully before answering" or "Work through this step by step"
Leverage system prompt	Put all persona, rules, constraints in system prompt via API — keep user messages clean
Long-context strength	Feed entire contracts, reports, or transcripts for deep analysis
Self-critique loop	Add: "Review your answer. What is the weakest part? How would you improve it?"

13.3 Claude System Prompt Template (API)

```
# SYSTEM PROMPT
You are [role] with [experience].
Your communication style is [tone].
Always [key rule 1] and [key rule 2].
Never [key restriction].

# USER MESSAGE
[Clear single instruction]
[Background, audience, goal]
[User content or question]
[Exact format required]
```